

# PyTorch\_Computer\_Vision\_MNIST\_Digits\_Classification

## 1 PyTorch Computer Vision MNIST Digits Classification

Author: Sufiyaan Usmani

Link: [GitHub](#)

## 2 Importing Libraries

```
[1]: import torch
      from torch import nn

      import torchvision
      from torchvision import datasets
      from torchvision.transforms import ToTensor

      import matplotlib.pyplot as plt
```

## 3 Setting up Device Agnostic Code

```
[2]: device = "cuda" if torch.cuda.is_available() else "cpu"
      print(device)
```

cpu

## 4 Getting a Dataset

Using MNIST Dataset

```
[3]: train_data = datasets.MNIST(
      root="data",
      train=True,
      download=True,
      transform=ToTensor(),
      target_transform=None
    )
```

```
test_data = datasets.MNIST(  
    root="data",  
    train=False,  
    download=True,  
    transform=ToTensor(),  
    target_transform=None  
)
```

```
[4]: image, label = train_data[0]  
     image.shape, label
```

```
[4]: (torch.Size([1, 28, 28]), 5)
```

```
[5]: len(train_data.data), len(train_data.targets), len(test_data.data),  
     ↪ len(test_data.targets)
```

```
[5]: (60000, 60000, 10000, 10000)
```


```
[6]: class_names = train_data.classes  
     class_names
```

```
[6]: ['0 - zero',  
      '1 - one',  
      '2 - two',  
      '3 - three',  
      '4 - four',  
      '5 - five',  
      '6 - six',  
      '7 - seven',  
      '8 - eight',  
      '9 - nine']
```


## 5 Visualizing Data

```
[7]: plt.imshow(image.squeeze(), cmap="gray")  
     plt.title(class_names[label])
```

```
[7]: Text(0.5, 1.0, '5 - five')
```



```
[8]: # Plot more images
fig = plt.figure(figsize=(9, 9))
rows, cols = 4, 4
for i in range(1, rows * cols + 1):
    random_idx = torch.randint(0, len(train_data), size=[1]).item()
    img, label = train_data[random_idx]
    fig.add_subplot(rows, cols, i)
    plt.imshow(img.squeeze(), cmap="gray")
    plt.title(class_names[label])
    plt.axis(False);
```



## 6 Preparing DataLoader

```
[9]: from torch.utils.data import DataLoader
```

```
[10]: BATCH_SIZE = 32

train_dataloader = DataLoader(
    dataset=train_data,
    batch_size=BATCH_SIZE,
    shuffle=True
)
```

```
test_dataloader = DataLoader(
    dataset=test_data,
    batch_size=BATCH_SIZE,
    shuffle=False
)
```

## 7 Building Model

Using TinyVGG

```
[11]: def TinyVGG(input_shape: int, hidden_units: int, output_shape: int):
    return nn.Sequential(
        nn.Conv2d(in_channels=input_shape, out_channels=hidden_units,
        ↪kernel_size=3, stride=1, padding=1),
        nn.ReLU(),
        nn.Conv2d(in_channels=hidden_units, out_channels=hidden_units,
        ↪kernel_size=3, stride=1, padding=1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2),

        nn.Conv2d(in_channels=hidden_units, out_channels=hidden_units,
        ↪kernel_size=3, stride=1, padding=1),
        nn.ReLU(),
        nn.Conv2d(in_channels=hidden_units, out_channels=hidden_units,
        ↪kernel_size=3, stride=1, padding=1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2),

        nn.Flatten(),
        nn.Linear(in_features=hidden_units*7*7, out_features=output_shape)
    )
```

```
[12]: model = TinyVGG(input_shape=1, hidden_units=10, output_shape=len(class_names)).
    ↪to(device)
model
```

```
[12]: Sequential(
  (0): Conv2d(1, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU()
  (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU()
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (5): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU()
```

```

(7): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): ReLU()
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(10): Flatten(start_dim=1, end_dim=-1)
(11): Linear(in_features=490, out_features=10, bias=True)
)

```

## 7.1 Train Step

```
[13]: from sklearn.metrics import accuracy_score
```

```
[14]: def train_step(model: torch.nn.Module, data_loader: torch.utils.data.
↳DataLoader, loss_fn: torch.nn.Module, optimizer: torch.optim.Optimizer,
↳accuracy_fn, device: torch.device=device):
    train_loss, train_acc = 0, 0
    model = model.to(device)
    model.train()
    for batch, (X, y) in enumerate(data_loader):
        X, y = X.to(device), y.to(device)
        y_pred = model(X)
        loss = loss_fn(y_pred, y)
        train_loss += loss
        train_acc += accuracy_fn(y_true=y, y_pred=y_pred.argmax(dim=1))

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    train_loss /= len(data_loader)
    train_acc /= len(data_loader)
    print(f"Train Loss: {train_loss:.5f} | Train Accuracy: {train_acc:.2f} %")

```

```
[15]: def test_step(model: torch.nn.Module, data_loader: torch.utils.data.DataLoader,
↳loss_fn: torch.nn.Module, accuracy_fn, device: torch.device=device):
    test_loss, test_acc = 0, 0
    model = model.to(device)
    model.eval()
    with torch.inference_mode():
        for X, y in data_loader:
            X, y = X.to(device), y.to(device)
            test_pred = model(X)
            test_loss += loss_fn(test_pred, y)
            test_acc += accuracy_fn(y_true=y, y_pred=test_pred.argmax(dim=1))

    test_loss /= len(data_loader)

```

```
test_acc /= len(data_loader)
print(f"Test Loss: {test_loss:.5f} | Test Accuracy: {test_acc:.2f} %")
```

## 8 Setting up loss function and optimizer

```
[16]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(params=model.parameters(), lr=0.1)
```

## 9 Train Test Loop

```
[17]: from tqdm.auto import tqdm
```

```
[18]: epochs = 10
```

```
[19]: for epoch in tqdm(range(epochs)):
    print(f"Epoch: {epoch}\n-----")
    train_step(model=model, data_loader=train_dataloader, loss_fn=loss_fn,
    ↪optimizer=optimizer, accuracy_fn=accuracy_score, device=device)
    test_step(model=model, data_loader=test_dataloader, loss_fn=loss_fn,
    ↪accuracy_fn=accuracy_score, device=device)
```

```
0%|          | 0/10 [00:00<?, ?it/s]
```

```
Epoch: 0
```

```
-----
```

```
Train Loss: 0.26906 | Train Accuracy: 0.91 %
```

```
Test Loss: 0.06224 | Test Accuracy: 0.98 %
```

```
Epoch: 1
```

```
-----
```

```
Train Loss: 0.06830 | Train Accuracy: 0.98 %
```

```
Test Loss: 0.06716 | Test Accuracy: 0.98 %
```

```
Epoch: 2
```

```
-----
```

```
Train Loss: 0.05411 | Train Accuracy: 0.98 %
```

```
Test Loss: 0.05236 | Test Accuracy: 0.98 %
```

```
Epoch: 3
```

```
-----
```

```
Train Loss: 0.04602 | Train Accuracy: 0.99 %
```

```
Test Loss: 0.04271 | Test Accuracy: 0.99 %
```

```
Epoch: 4
```

```
-----
```

```
Train Loss: 0.04066 | Train Accuracy: 0.99 %
```

```
Test Loss: 0.04345 | Test Accuracy: 0.99 %
```

```
Epoch: 5
```

```
-----
```

```
Train Loss: 0.03751 | Train Accuracy: 0.99 %
```

```
Test Loss: 0.03820 | Test Accuracy: 0.99 %  
Epoch: 6  
-----  
Train Loss: 0.03315 | Train Accuracy: 0.99 %  
Test Loss: 0.03552 | Test Accuracy: 0.99 %  
Epoch: 7  
-----  
Train Loss: 0.03092 | Train Accuracy: 0.99 %  
Test Loss: 0.03731 | Test Accuracy: 0.99 %  
Epoch: 8  
-----  
Train Loss: 0.02830 | Train Accuracy: 0.99 %  
Test Loss: 0.03752 | Test Accuracy: 0.99 %  
Epoch: 9  
-----  
Train Loss: 0.02675 | Train Accuracy: 0.99 %  
Test Loss: 0.03710 | Test Accuracy: 0.99 %
```

## 10 Making a prediction

```
[26]: batch, (X, y) = next(enumerate((test_dataloader)))
```

```
[41]: X[1].shape, y[1]
```


```
[41]: (torch.Size([1, 28, 28]), tensor(2))
```

```
[28]: img, label = X[1], y[1]
```

```
[29]: plt.imshow(img.squeeze(), cmap="gray")  
      plt.title(class_names[label])
```

```
[29]: Text(0.5, 1.0, '2 - two')
```





```
[33]: model.eval()
      with torch.inference_mode():
          y_pred = model(X)
```

```
[45]: pred_label = y_pred.argmax(dim=1)[1]
```

```
[46]: print(f"Prediction: {class_names[pred_label]}")
```

Prediction: 2 - two

## 11 Confusion Matrix

```
[47]: # Import tqdm for progress bar
      from tqdm.auto import tqdm

      # 1. Make predictions with trained model
      y_preds = []
      model.eval()
      with torch.inference_mode():
          for X, y in tqdm(test_dataloader, desc="Making predictions"):
```

```

# Send data and targets to target device
X, y = X.to(device), y.to(device)
# Do the forward pass
y_logit = model(X)
# Turn predictions from logits -> prediction probabilities -> predictions
↳ labels
y_pred = torch.softmax(y_logit, dim=1).argmax(dim=1) # note: perform
↳ softmax on the "logits" dimension, not "batch" dimension (in this case we
↳ have a batch size of 32, so can perform on dim=1)
# Put predictions on CPU for evaluation
y_preds.append(y_pred.cpu())
# Concatenate list of predictions into a tensor
y_pred_tensor = torch.cat(y_preds)

```

Making predictions: 0%| | 0/313 [00:00<?, ?it/s]

```

[48]: # See if torchmetrics exists, if not, install it
try:
    import torchmetrics, mlxtend
    print(f"mlxtend version: {mlxtend.__version__}")
    assert int(mlxtend.__version__.split(".")[1]) >= 19, "mlxtend verison
↳ should be 0.19.0 or higher"
except:
    !pip install -q torchmetrics -U mlxtend # <- Note: If you're using Google
↳ Colab, this may require restarting the runtime
    import torchmetrics, mlxtend
    print(f"mlxtend version: {mlxtend.__version__}")

```

840.4/840.4

kB 7.4 MB/s eta 0:00:00

1.4/1.4 MB

12.5 MB/s eta 0:00:00

mlxtend version: 0.23.1

```


[49]: from torchmetrics import ConfusionMatrix
from mlxtend.plotting import plot_confusion_matrix

# 2. Setup confusion matrix instance and compare predictions to targets
confmat = ConfusionMatrix(num_classes=len(class_names), task='multiclass')
confmat_tensor = confmat(preds=y_pred_tensor,
                        target=test_data.targets)

# 3. Plot the confusion matrix
fig, ax = plot_confusion_matrix(
    conf_mat=confmat_tensor.numpy(), # matplotlib likes working with NumPy
    class_names=class_names, # turn the row and column labels into class names

```

```
figsize=(10, 7)
);
```



## 12 Saving and Loading Model

```
[50]: from pathlib import Path

# Create models directory (if it doesn't already exist), see: https://docs.
# python.org/3/library/pathlib.html#pathlib.Path.mkdir
MODEL_PATH = Path("models")
MODEL_PATH.mkdir(parents=True, # create parent directories if needed
                  exist_ok=True # if models directory already exists, don't error
```

```

)

# Create model save path
MODEL_NAME = "mnist_cnn_model.pth"
MODEL_SAVE_PATH = MODEL_PATH / MODEL_NAME

# Save the model state dict
print(f"Saving model to: {MODEL_SAVE_PATH}")
torch.save(obj=model.state_dict(), # only saving the state_dict() only saves_
    ↪ the learned parameters
           f=MODEL_SAVE_PATH)

```

Saving model to: models/mnist\_cnn\_model.pth

```

[51]: loaded_model = TinyVGG(input_shape=1,
                             hidden_units=10,
                             output_shape=10)

# Load in the saved state_dict()
loaded_model.load_state_dict(torch.load(f=MODEL_SAVE_PATH))

loaded_model = loaded_model.to(device)

```

```

[52]: loaded_model

```

```

[52]: Sequential(
  (0): Conv2d(1, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU()
  (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU()
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (5): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (6): ReLU()
  (7): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (8): ReLU()
  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (10): Flatten(start_dim=1, end_dim=-1)
  (11): Linear(in_features=490, out_features=10, bias=True)
)

```

```

[53]: model.eval()
with torch.inference_mode():
    y_pred = loaded_model(X)

pred_label = y_pred.argmax(dim=1)[1]

```

```
print(f"Prediction using Loaded Model: {class_names[pred_label]}")
```

Prediction using Loaded Model: 2 - two